

Comprehensive Presentation & Technical Understanding Guide

Prepared by Zarar | AiTeC Intelligence Systems Group

This expanded guide is designed to give you a deep, confident understanding of the system's inner workings so you can easily handle technical questions from your supervisor.

Part 1: The Core Philosophy (Algorithms vs. Models)

One of the most important design decisions in TrendChecker is why we used Mathematical Algorithms instead of Pre-trained Neural Network Models (Deep Learning).

If your supervisor asks, "Why didn't you just use a pre-trained deep learning model like ResNet or a HuggingFace AI detector?", here is how you answer:

The Difference:

- **Pre-Trained Models (Deep Learning):** These are trained on massive datasets to recognize patterns (e.g., feeding an AI 10,000 real faces and 10,000 AI faces). **The Problem:** They are "Black Boxes" (you don't know why they made a decision). They require heavy GPUs, Gigabytes of storage, and they suffer from "Concept Drift" (if Midjourney releases a new update, the pre-trained model becomes outdated and must be retrained).
- **Deterministic Algorithms (Our Approach):** These are hard-coded mathematical formulas that analyze absolute physical or structural truths. **The Solution:** Algorithms run instantly on a basic CPU, take up almost zero disk space, and are 100% interpretable (we can point to the exact math that flagged a fake). Most importantly, they don't care how "good" an AI generator gets. Even the best AI in the world does not have a physical camera lens, so mathematical forensics will always catch it.

Conclusion for Supervisor: By using mathematical algorithms, TrendChecker achieves 100% local processing, zero GPU reliance, complete data privacy, and immunity to AI generator updates.

Part 2: Detailed Pipeline Breakdown

1. Text Analysis Pipeline (Lexical Shannon Entropy)

The Basic Pitch: We analyze the text to see if a human wrote it or if an AI (like ChatGPT) generated it.

The Technical Deep Dive:

How it works: Large Language Models (LLMs) operate on probability. When they write a sentence, they are just predicting the most mathematically likely next word. Because of this, AI text is highly predictable.

The Algorithm: We use Shannon Entropy, a concept from Information Theory that measures randomness.

For your understanding: Imagine rolling a loaded die that always lands on 6 (AI). The entropy (randomness) is very low. Now imagine rolling a fair die (Human). The entropy is high. We calculate the probability distribution of the words in a tweet. If the words are too predictable, the entropy score drops, and we flag it as AI_GENERATED. Humans make typos, use weird slang, and structure sentences chaotically—which creates high entropy.

2. Claim Fact-Checking Pipeline (Domain Exclusion)

The Basic Pitch: We extract the core claim of a post and verify it against trusted news sources while preventing "echo chambers."

The Technical Deep Dive:

How it works: Misinformation spreads because fake news sites cite each other. If we just search Google, we might find 10 websites confirming a fake story, but they are all linking back to the original fake tweet.

The Algorithm: We use Entity Extraction followed by Domain Exclusion Search.

For your understanding: If a claim originates from *reddit.com/r/news*, our algorithm actively blocks *reddit.com* from the search results. It forces the system to look outward. We then require a Multi-Source Consensus—meaning the system must find the entities mentioned in the claim on 3 to 5 independent, whitelisted domains (like *reuters.com* or *apnews.com*). If it only finds it on unverified blogs, the claim is marked UNSUPPORTED.

3. Image Forensics Pipeline (2D FFT & PRNU Kurtosis)

The Basic Pitch: We detect AI-generated images by looking at the invisible "noise" in the photo, rather than relying on easily-deleted metadata.

The Technical Deep Dive:

How it works: Every physical camera sensor (even in an iPhone) has microscopic manufacturing defects. When a photo is taken, these defects leave a unique, invisible fingerprint of noise called PRNU (Photo Response Non-Uniformity).

The Algorithm: AI image generators (like Midjourney) create images purely from math—they have no physical sensor, so they lack this natural PRNU noise. Their pixels are unnaturally smooth.

For your understanding: We use a 2D Fast Fourier Transform (FFT). FFT translates an image from the "spatial domain" (pixels you see) into the "frequency domain" (waves). In the frequency domain, we calculate the Kurtosis (the sharpness of the data peaks).

- A real photo has random, scattered noise (Low Kurtosis).
- An AI photo has repeating, mathematical artifact patterns (High Kurtosis). We measure that spike to catch the AI.

4. Video Forensics Pipeline (Gunnar Farneback Optical Flow)

***The Basic Pitch:** We detect deepfakes and AI-generated video (like OpenAI Sora) by analyzing how objects move across the screen to see if they obey the laws of physics.*

The Technical Deep Dive:

How it works: AI video generators don't actually film moving objects; they hallucinate a sequence of images. Because of this, they struggle with "temporal consistency" (keeping objects perfectly solid as they move over time).

The Algorithm: We extract keyframes directly into RAM and apply the Gunnar Farneback Optical Flow algorithm.

For your understanding: Optical Flow assigns a "motion vector" (speed and direction) to every pixel in a frame. If a person waves their hand, the pixels should move together in a smooth, continuous vector. In AI-generated video, pixels often "teleport" or morph randomly between frames because the AI is guessing. Our algorithm tracks these vectors. If the vectors show impossible physical movement or sudden chaotic morphing, the video is flagged as LIKELY_AI_VIDEO.

5. The Authenticity Engine (Composite Scoring)

***The Wrap-Up:** Finally, we don't treat these pipelines in isolation.*

The Authenticity Engine gathers the Lexical Entropy score, the Fact-Check consensus count, the Image Kurtosis value, and the Video Motion Vectors. It applies a weighted mathematical formula to generate a single, easy-to-understand Confidence Score (0-100%).

This proves to your supervisor that the system isn't just a collection of random scripts, but a cohesive, engineered forensic engine.